

Aprendizado Supervisionado

Regressão Logística

1 - Bibliotecas

```
In [1]: %matplotlib notebook
import numpy as np
import pandas as pd
import seaborn as sn
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn import neighbors
from sklearn.datasets import make_blobs

from matplotlib.colors import ListedColormap

cmap_light = ListedColormap(['#FFFFAA', '#AAFFAA', '#AAAAFF', '#EEEEFF'])
cmap_bold = ListedColormap(['#FFFF00', '#00FF00', '#0000FF', '#000000'])
```

2 - Funções para plotagem

```
In [2]: def plot_class_regions_for_classifier_subplot(clf, X, y,
                                                    X_test, y_test,
                                                    title, subplot,
                                                    target_names = None,
                                                    plot_decision_regions = True):

    numClasses = np.amax(y) + 1
    color_list_light = ['#FFFFAA', '#EEEEFF', '#AAFFAA', '#AAAAFF']
    color_list_bold = ['#EEEE00', '#000000', '#00CC00', '#0000CC']
    cmap_light = ListedColormap(color_list_light[0:numClasses])
    cmap_bold = ListedColormap(color_list_bold[0:numClasses])

    h = 0.03
    k = 0.5
    x_plot_adjust = 0.1
    y_plot_adjust = 0.1
    plot_symbol_size = 50

    x_min = X[:, 0].min()
    x_max = X[:, 0].max()
    y_min = X[:, 1].min()
    y_max = X[:, 1].max()
    x2, y2 = np.meshgrid(np.arange(x_min-k, x_max+k, h), np.arange(y_min-k, y_max+k, h))

    P = clf.predict(np.c_[x2.ravel(), y2.ravel()])
    P = P.reshape(x2.shape)

    if plot_decision_regions:
        subplot.contourf(x2, y2, P, cmap=cmap_light, alpha = 0.8)

    subplot.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap_bold, s=plot_symbol_size, edgecolor =
    subplot.set_xlim(x_min - x_plot_adjust, x_max + x_plot_adjust)
    subplot.set_ylim(y_min - y_plot_adjust, y_max + y_plot_adjust)
```

```

if (X_test is not None):
    subplot.scatter(X_test[:, 0], X_test[:, 1], c=y_test, cmap=cmap_bold, s=plot_symbol_size)
    train_score = clf.score(X, y)
    test_score = clf.score(X_test, y_test)
    title = title + "\nTreinamento = {:.2f}, Teste = {:.2f}".format(train_score, test_score)

subplot.set_title(title)

if (target_names is not None):
    legend_handles = []
    for i in range(0, len(target_names)):
        patch = mpatches.Patch(color=color_list_bold[i], label=target_names[i])
        legend_handles.append(patch)
    subplot.legend(loc=0, handles=legend_handles)

def plot_class_regions_for_classifier(clf, X, y,
                                     X_test=None, y_test=None,
                                     title=None,
                                     target_names = None,
                                     plot_decision_regions = True):

    numClasses = np.amax(y) + 1
    color_list_light = ['#FFFFFF', '#E0E0E0', '#A0A0A0', '#AAAAFF']
    color_list_bold = ['#000000', '#000000', '#00CC00', '#0000CC']
    cmap_light = ListedColormap(color_list_light[0:numClasses])
    cmap_bold = ListedColormap(color_list_bold[0:numClasses])

    h = 0.03
    k = 0.5
    x_plot_adjust = 0.1
    y_plot_adjust = 0.1
    plot_symbol_size = 50

    x_min = X[:, 0].min()
    x_max = X[:, 0].max()
    y_min = X[:, 1].min()
    y_max = X[:, 1].max()
    x2, y2 = np.meshgrid(np.arange(x_min-k, x_max+k, h), np.arange(y_min-k, y_max+k, h))

    P = clf.predict(np.c_[x2.ravel(), y2.ravel()])
    P = P.reshape(x2.shape)
    plt.figure()
    if plot_decision_regions:
        plt.contourf(x2, y2, P, cmap=cmap_light, alpha = 0.8)

    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap_bold, s=plot_symbol_size, edgecolor = 'b')
    plt.xlim(x_min - x_plot_adjust, x_max + x_plot_adjust)
    plt.ylim(y_min - y_plot_adjust, y_max + y_plot_adjust)

    if (X_test is not None):
        plt.scatter(X_test[:, 0], X_test[:, 1], c=y_test, cmap=cmap_bold, s=plot_symbol_size)
        train_score = clf.score(X, y)
        test_score = clf.score(X_test, y_test)
        title = title + "\nTrain score = {:.2f}, Test score = {:.2f}".format(train_score, test_score)

    if (target_names is not None):
        legend_handles = []
        for i in range(0, len(target_names)):
            patch = mpatches.Patch(color=color_list_bold[i], label=target_names[i])
            legend_handles.append(patch)
        plt.legend(loc=0, handles=legend_handles)

    if (title is not None):
        plt.title(title)
    plt.show()

```

3 - Exemplo de regressão logística

In [5]:

```
from sklearn.linear_model import LogisticRegression

fruits = pd.read_table('./Data/fruit_data_with_colors.txt')

#fruits.head()

X_fruits_2d = fruits[['height', 'width']]
y_fruits_2d = fruits['fruit_label']

fig, subaxes = plt.subplots(1, 1, figsize=(7, 5))
y_fruits_apple = y_fruits_2d == 1  # apples vs todas as outras frutas

X_train, X_test, y_train, y_test = (train_test_split(X_fruits_2d.to_numpy(),
                                                    y_fruits_apple.to_numpy(),
                                                    random_state = 0))

clf = LogisticRegression(C=100).fit(X_train, y_train)


plot_class_regions_for_classifier_subplot(clf, X_train, y_train, None,
                                         None, 'Apple vs outras frutas',
                                         subaxes)

print(clf.predict([[6,8]]))

print(clf.predict([[10,7]]))

print(clf.score(X_train, y_train))

print(clf.score(X_test, y_test))
```



```
[ True]
[False]
0.7954545454545454
0.7333333333333333
```

4 - Criação de dados sintéticos

In [6]:


```
from sklearn.datasets import make_classification

plt.figure()

X_C2, y_C2 = make_classification(n_samples = 100, n_features=2,
                                n_redundant=0, n_informative=2,
                                n_clusters_per_class=1, flip_y = 0.2,
                                class_sep = 0.5, random_state=0)

plt.scatter(X_C2[:, 0], X_C2[:, 1], c=y_C2, marker='o', cmap=cmap_bold)

plt.show()
```



5 - Regressão logística nos dados sintéticos

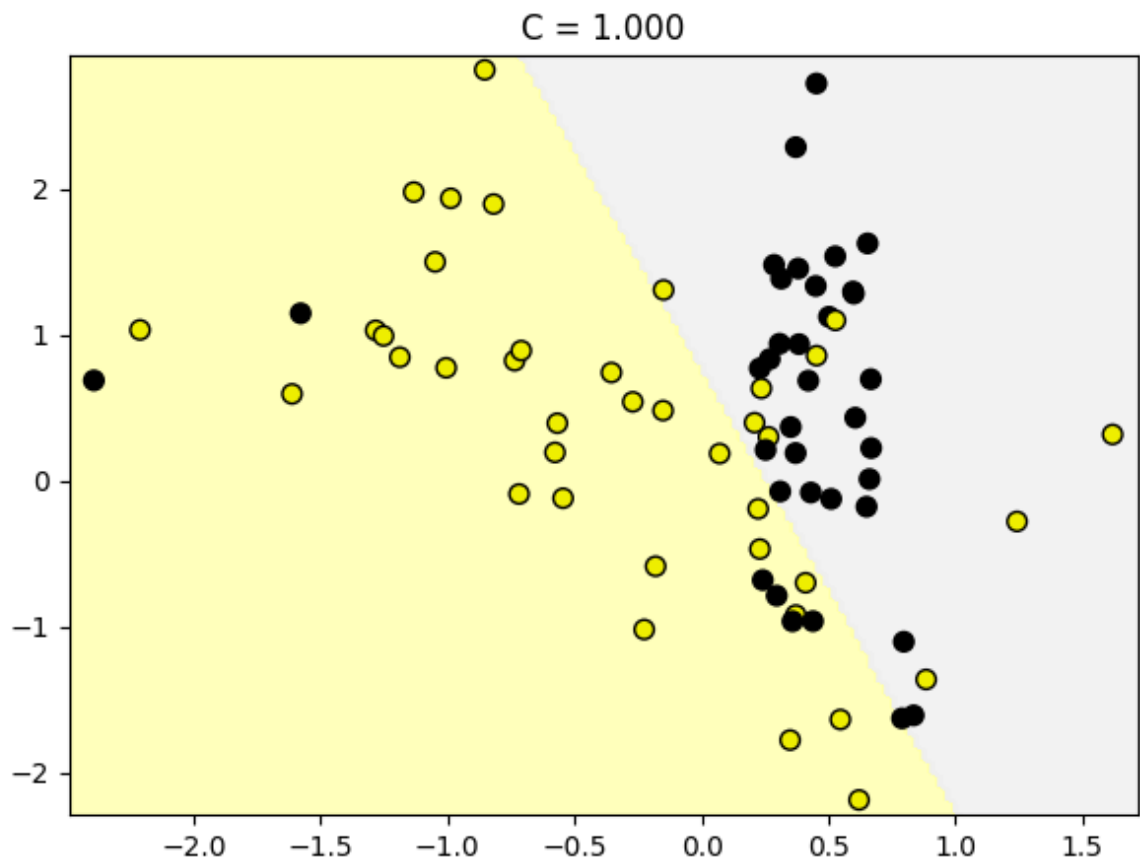
In [7]:

```
from sklearn.linear_model import LogisticRegression

X_train, X_test, y_train, y_test = train_test_split(X_C2, y_C2,
                                                    random_state = 0)

fig, subaxes = plt.subplots(1, 1, figsize=(7, 5))
clf = LogisticRegression().fit(X_train, y_train)
title = 'C = {:.3f}'.format(1.0)
plot_class_regions_for_classifier_subplot(clf, X_train, y_train,
                                          None, None, title, subaxes)

print(clf.score(X_train, y_train))
print(clf.score(X_test, y_test))
```



0.786666666666666666


0.76

5 - Regressão logística e o parâmetro C


In [9]:

[illegible]


Apple vs. outras frutas, $C = 0.100$
Treinamento = 0.66, Teste = 0.67



Apple vs. outras frutas, $C = 1.000$
Treinamento = 0.75, Teste = 0.67



Apple vs. outras frutas, $C = 100.000$
Treinamento = 0.80, Teste = 0.73



6 - O exemplo linear

In [10]:

```

from sklearn.svm import SVC

X_train, X_test, y_train, y_test = train_test_split(X_C2, y_C2, random_state = 0)


fig, subaxes = plt.subplots(1, 1, figsize=(7, 5))
this_C = 1.0

clf = SVC(kernel = 'linear', C=this_C).fit(X_train, y_train)

title = 'SVC linear, C = {:.3f}'.format(this_C)
plot_class_regions_for_classifier_subplot(clf, X_train, y_train, None, None, title, subaxes)

print(clf.score(X_train, y_train))
print(clf.score(X_test, y_test))

```



0.76

0.8

7 - SVM e o parâmetro C

In [11]:

```

from sklearn.svm import LinearSVC

X_train, X_test, y_train, y_test = train_test_split(X_C2, y_C2, random_state = 0)

fig, subaxes = plt.subplots(1, 2, figsize=(8, 4))

for this_C, subplot in zip([0.00001, 100], subaxes):
    clf = LinearSVC(C=this_C).fit(X_train, y_train)

```




```

title = 'SVC linear, C = {:.5f}'.format(this_C)
print(clf.score(X_train, y_train))
print(clf.score(X_test, y_test))
print('\n')
plot_class_regions_for_classifier_subplot(clf, X_train, y_train,
                                         None, None, title, subplot)

plt.tight_layout()

```



```

0.7733333333333333
0.8

```

```

0.7466666666666667
0.8

```

```

/Users/marinaramalhetedesouza/opt/anaconda3/envs/ml-impa/lib/python3.8/site-packages/sklearn/svm/_base.py:1206: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
  warnings.warn(

```

Máquinas de Vetores de Suportes - Kernelização

(Kernelized Support Vector Machines)

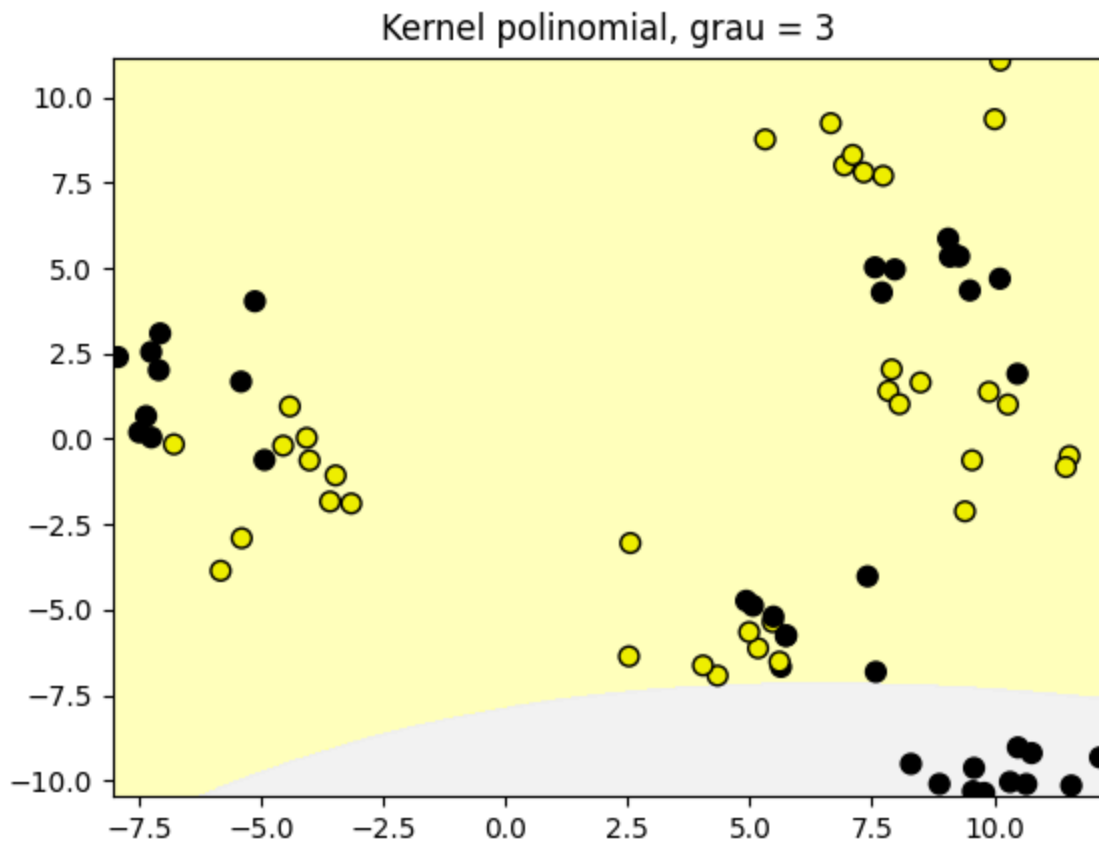
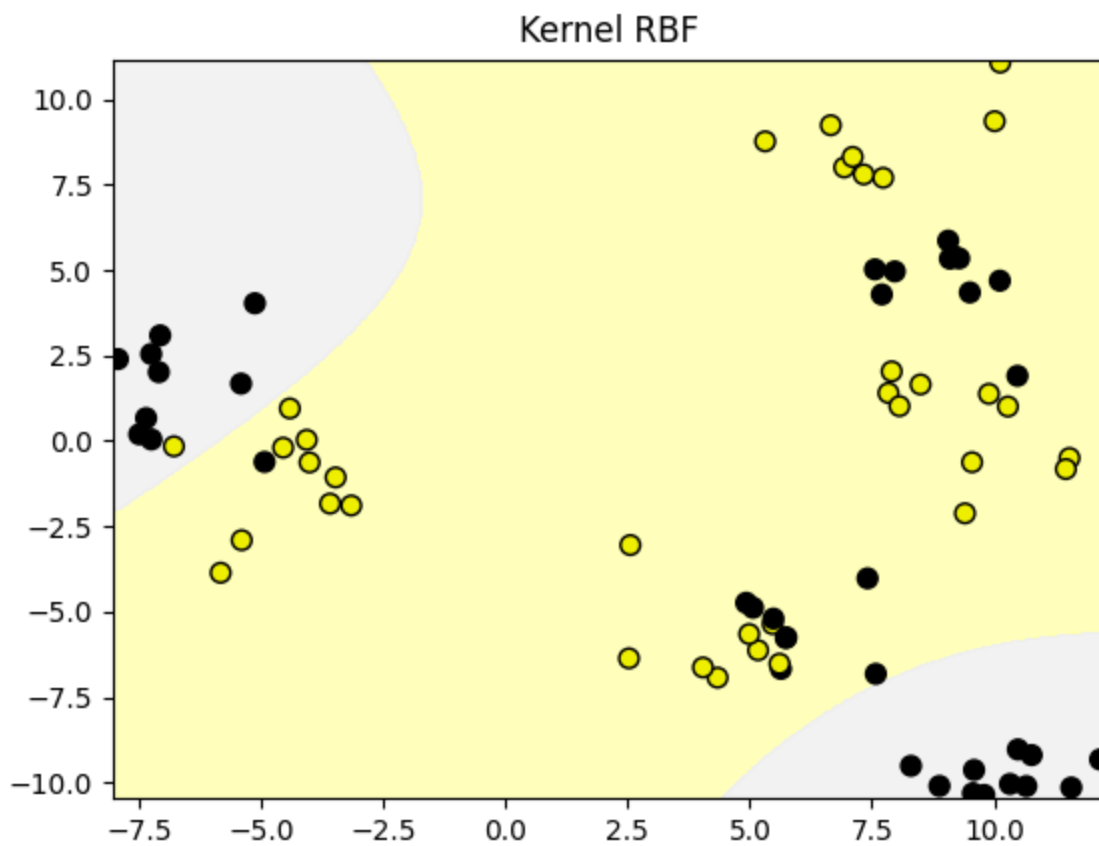
8 - Base de dados sintética com problema complexo de classificação

```

In [12]: X_D2, y_D2 = make_blobs(n_samples = 100, n_features = 2, centers = 8,
                                cluster_std = 1.3, random_state = 4)

y_D2 = y_D2 % 2
plt.figure()
plt.title('Problema complexo de classificação')
plt.scatter(X_D2[:,0], X_D2[:,1], c=y_D2,
            marker='o', s=50, cmap=cmap_bold)
plt.show()

```


```
X_train, X_test, y_train, y_test = train_test_split(X_D2, y_D2, random_state = 0)

fig, subaxes = plt.subplots(3, 1, figsize=(4, 11))


for this_gamma, subplot in zip([0.01, 1.0, 10.0], subaxes):
    clf = SVC(kernel = 'rbf', gamma=this_gamma).fit(X_train, y_train)
    title = 'Gamma = {:.2f}'.format(this_gamma)
    plot_class_regions_for_classifier_subplot(clf, X_train, y_train,
                                             None, None, title, subplot)

plt.tight_layout()
```


Gamma = 0.01



Gamma = 1.00



Gamma = 10.00



11 - Kernel RBF + parâmetros C e Gamma

In [15]:

```
from sklearn.svm import SVC


X_train, X_test, y_train, y_test = train_test_split(X_D2, y_D2, random_state = 0)

fig, subaxes = plt.subplots(3, 4, figsize=(15, 10), dpi=50)

for this_gamma, this_axis in zip([0.01, 1, 5], subaxes):

    for this_C, subplot in zip([0.1, 1, 15, 250], this_axis):
        title = 'gamma = {:.2f}, C = {:.2f}'.format(this_gamma, this_C)
        clf = SVC(kernel = 'rbf', gamma = this_gamma,
                   C = this_C).fit(X_train, y_train)
        plot_class_regions_for_classifier_subplot(clf, X_train, y_train,
                                                X_test, y_test, title,
                                                subplot)

plt.tight_layout(pad=0.4, w_pad=0.5, h_pad=1.0)
```



In []: