

# Avaliação dos modelos

## 1 - Bibliotecas

```
In [1]: %matplotlib notebook
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn import neighbors

import graphviz
from sklearn.tree import export_graphviz
```

## 2 - Carregando a base de dados sobre frutas

```
In [2]: fruits = pd.read_table('./Data/fruit_data_with_colors.txt')

X_fruits_2d = fruits[['height', 'width']]
y_fruits_2d = fruits['fruit_label']

fruits.head()
```

```
Out[2]:
```

	fruit_label	fruit_name	fruit_subtype	mass	width	height	color_score
0	1	apple	granny_smith	192	8.4	7.3	0.55
1	1	apple	granny_smith	180	8.0	6.8	0.59
2	1	apple	granny_smith	176	7.4	7.2	0.60
3	2	mandarin	mandarin	86	6.2	4.7	0.80
4	2	mandarin	mandarin	84	6.0	4.6	0.79

```
In [4]: fruits.shape
```

```
Out[4]: (59, 7)
```

# Validação Cruzada

## 3 - Validação Cruzada

```
In [5]: from sklearn.model_selection import cross_val_score

clf = KNeighborsClassifier(n_neighbors = 5)
X = X_fruits_2d.values
y = y_fruits_2d.values
cv_scores = cross_val_score(clf, X, y, cv=5)

print('Validação cruzada:', cv_scores)
print('Média: {:.3f}'.format(np.mean(cv_scores)))
```

```
Validação cruzada: [0.75      0.75      0.83333333 0.83333333 0.81818182]
Média: 0.797
```

## 4 - Curva de Validação

In [6]:

```
from sklearn.svm import SVC
from sklearn.model_selection import validation_curve

param_range = np.logspace(-3, 3, 4)
train_scores, test_scores = validation_curve(SVC(C=1), X, y,
                                             param_name='gamma',
                                             param_range=param_range, cv=5)

print(train_scores)
print(test_scores)
```

```
[[0.46808511 0.40425532 0.40425532 0.34042553 0.33333333]
 [0.82978723 0.78723404 0.76595745 0.74468085 0.75       ]
 [0.87234043 0.89361702 0.89361702 0.89361702 0.85416667]
 [0.9787234  1.         0.9787234  1.         0.97916667]]
[[0.58333333 0.33333333 0.33333333 0.25       0.27272727]
 [0.83333333 0.66666667 0.66666667 0.75       0.72727273]
 [0.41666667 0.66666667 0.83333333 0.83333333 0.81818182]
 [0.33333333 0.33333333 0.25       0.33333333 0.36363636]]
```

## 5 - Plot da Validação

In [7]:

```
# Exemplo do scikit-learn: validation_plot
# http://scikit-learn.org/stable/auto\_examples/model\_selection/plot\_validation\_curve.html
plt.figure()

train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)

plt.title('Validation Curve with SVM')
plt.xlabel('$\gamma$ (gamma)')
plt.ylabel('Score')
plt.ylim(0.0, 1.1)
lw = 2


plt.semilogx(param_range, train_scores_mean, label='Escore de Treinamento',
              color='darkorange', lw=lw)

plt.fill_between(param_range, train_scores_mean - train_scores_std,
                 train_scores_mean + train_scores_std, alpha=0.2,
                 color='darkorange', lw=lw)

plt.semilogx(param_range, test_scores_mean, label='Escore da validação cruzada',
              color='navy', lw=lw)

plt.fill_between(param_range, test_scores_mean - test_scores_std,
                 test_scores_mean + test_scores_std, alpha=0.2,
                 color='navy', lw=lw)

plt.legend(loc='best')
plt.show()
```



## Árvores de Decisão

### 6 - Importar dataset Íris e executar árvore de decisão

```
In [8]: from sklearn.datasets import load_iris
from sklearn.tree import DecisionTreeClassifier

iris = load_iris()

X_train, X_test, y_train, y_test = train_test_split(iris.data, iris.target, random_state = 42)

clf = DecisionTreeClassifier().fit(X_train, y_train)

print('Acurácia da árvore de decisão no conjunto de treinamento: {:.2f}'.format(clf.score(X_train, y_train)))
print('Acurácia da árvore de decisão no conjunto de teste: {:.2f}'.format(clf.score(X_test, y_test)))
```

Acurácia da árvore de decisão no conjunto de treinamento: 1.00  
Acurácia da árvore de decisão no conjunto de teste: 0.95

### 7 - Profundidade da árvore (evitar overfitting)

```
In [9]: clf2 = DecisionTreeClassifier(min_samples_leaf = 10).fit(X_train, y_train)

print('Acurácia da árvore de decisão no conjunto de treinamento: {:.2f}'.format(clf2.score(X_train, y_train)))
print('Acurácia da árvore de decisão no conjunto de teste: {:.2f}'.format(clf2.score(X_test, y_test)))
```

Acurácia da árvore de decisão no conjunto de treinamento: 0.96  
Acurácia da árvore de decisão no conjunto de teste: 0.95

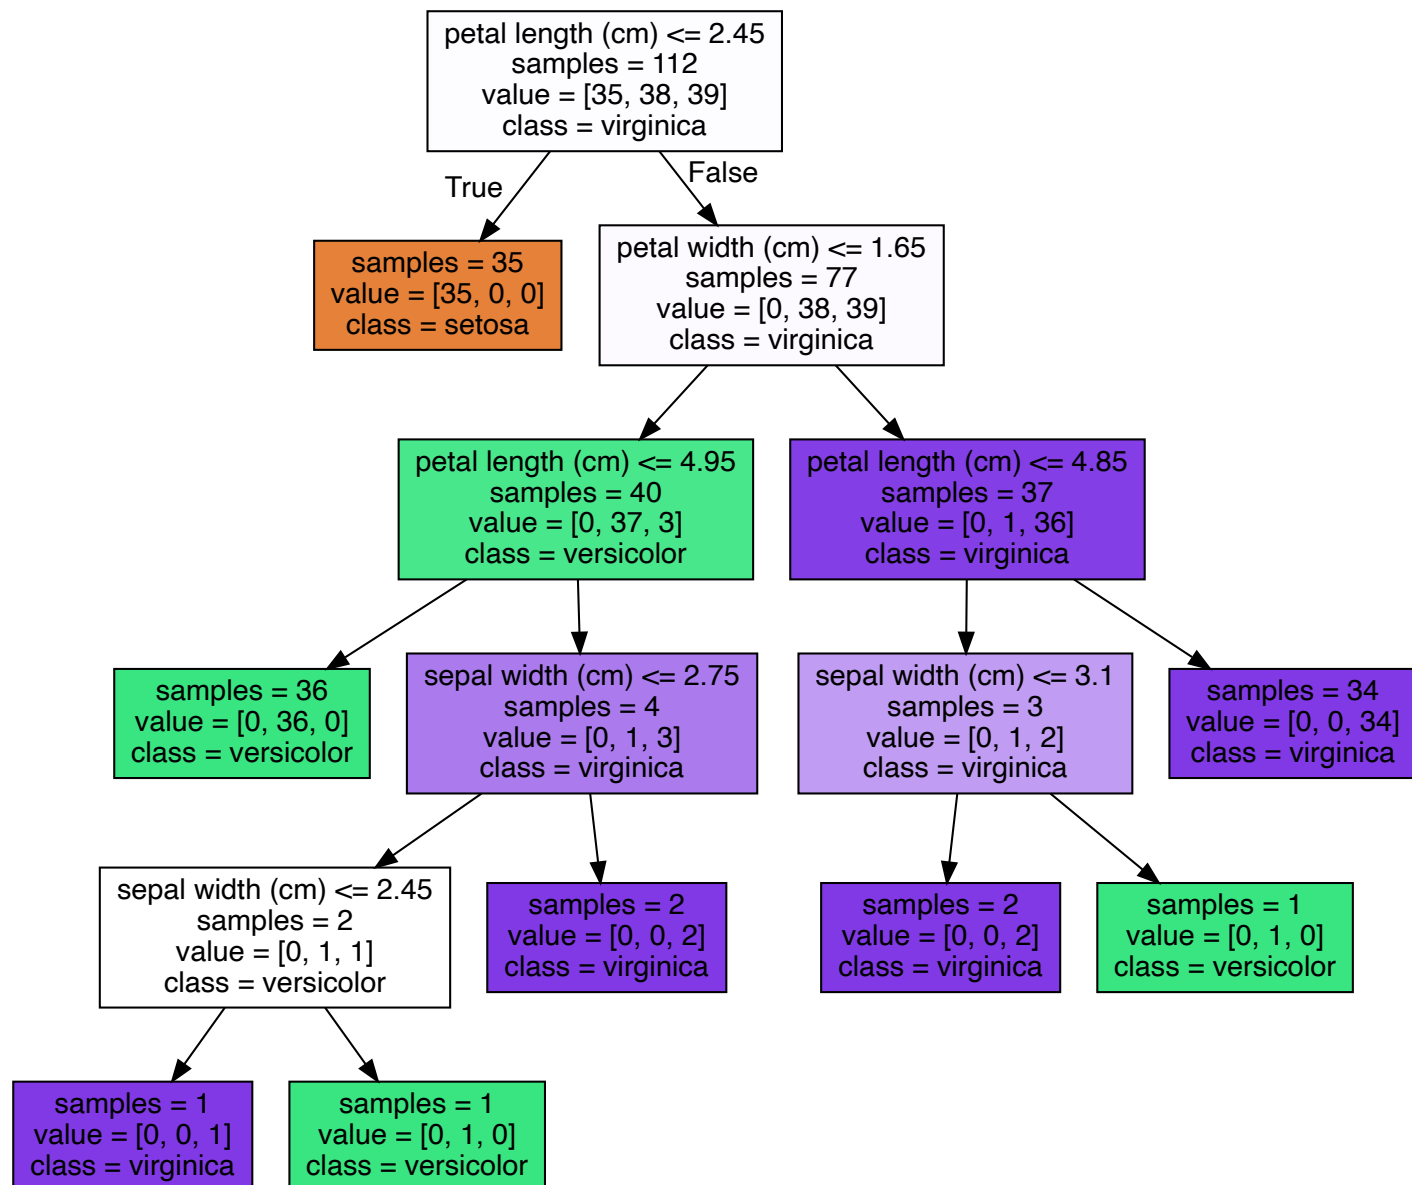
```
In [31]: # conda install python-graphviz
```

## 8 - Visualizando árvores de decisão

In [15]:

```
def plot_decision_tree(clf, feature_names, class_names):  
    export_graphviz(clf, out_file="adspy_temp.dot", feature_names=feature_names,  
                    class_names=class_names, filled = True, impurity = False)  
    with open("adspy_temp.dot") as f:  
        dot_graph = f.read()  
    return graphviz.Source(dot_graph)  
  
plot_decision_tree(clf, iris.feature_names, iris.target_names)
```

Out[15]:

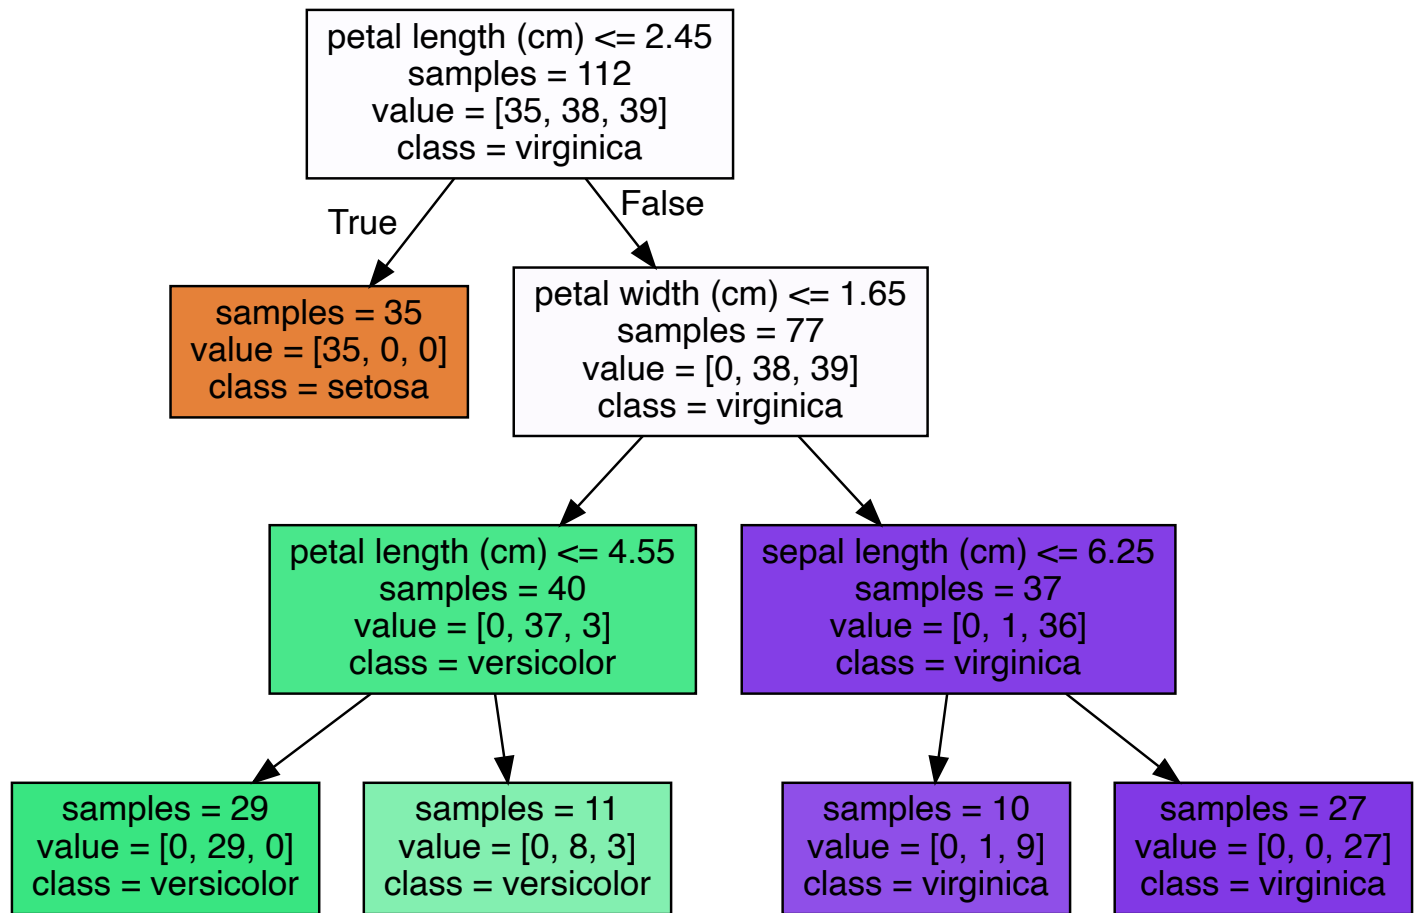


## 9 - Pré-podagem

In [16]:

```
plot_decision_tree(clf2, iris.feature_names, iris.target_names)
```

Out[16]:




## 10 - Importância da Característica

```
In [17]: def plot_feature_importances(clf, feature_names):
c_features = len(feature_names)
plt.barh(range(c_features), clf.feature_importances_)
plt.xlabel("Importância")
plt.ylabel("Característica")
plt.yticks(np.arange(c_features), feature_names)
```

```
In [18]: plt.figure(figsize=(10,4), dpi=80)
plot_feature_importances(clf2, iris.feature_names)
plt.show()

print('Importâncias: {}'.format(clf2.feature_importances_))
```



# Classificadores dummy e Base de dados pré-carregadas

## 11 - Carregando base de dados (digits)

```
In [19]: from sklearn.datasets import load_digits

dataset = load_digits()

X, y = dataset.data, dataset.target

for class_name, class_count in zip(dataset.target_names, np.bincount(dataset.target)):
    print(class_name, class_count)
```

0 178  
1 182  
2 177  
3 183  
4 181  
5 182  
6 181  
7 179  
8 174  
9 180

## 12 - Transformando a base de dados para não-balanceada

```
In [20]: y_binary_imbalanced = y.copy()
y_binary_imbalanced[y_binary_imbalanced != 1] = 0

print('Original:\t', y[1:30])
print('Novo:\t', y_binary_imbalanced[1:30])
```

Original: [1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9]  
Novo: [1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0]

## 13 - Verificando proporções da base de dados não-balanceada

```
In [21]: np.bincount(y_binary_imbalanced)
```

Out[21]: array([1615, 182])

## 14 - Treinando um classificador SVC

```
In [22]: X_train, X_test, y_train, y_test = train_test_split(X, y_binary_imbalanced, random_state=0)

from sklearn.svm import SVC

svm = SVC(kernel='rbf', C=1).fit(X_train, y_train)
svm.score(X_test, y_test)
```

Out[22]: 0.9955555555555555

## 15 - Classificadores "bôbos" Dummy

```
In [23]: from sklearn.dummy import DummyClassifier

# DummyClassifier?
```

[illegible]

```
In [24]: dummy_majority.score(X_test, y_test)
```

```
Out[24]: 0.9044444444444445
```

```
In [25]: svm = SVC(kernel='linear', C=1).fit(X_train, y_train)
svm.score(X_test, y_test)
```

```
Out[25]: 0.9777777777777777
```

```
In [26]: from sklearn.metrics import confusion_matrix

dummy_majority = DummyClassifier(strategy = 'most_frequent').fit(X_train, y_train)
y_majority_predicted = dummy_majority.predict(X_test)
confusion = confusion_matrix(y_test, y_majority_predicted)

print('Classe mais frequente\n', confusion)
```

```
Classe mais frequente
[[407    0]
 [ 43    0]]
```

## 19 - Matriz de confusão - classificador Dummy

```
In [27]: dummy_classprop = DummyClassifier(strategy='stratified').fit(X_train, y_train)
y_classprop_predicted = dummy_classprop.predict(X_test)
confusion = confusion_matrix(y_test, y_classprop_predicted)

print('Estratificado\n', confusion)
```

```
Estratificado
[[370  37]
 [ 39   4]]
```

## 20 - Matriz de confusão - SVC

```
In [28]: svm = SVC(kernel='linear', C=1).fit(X_train, y_train)
svm_predicted = svm.predict(X_test)
confusion = confusion_matrix(y_test, svm_predicted)

print('SVC (kernel linear, C=1)\n', confusion)
```

```
SVC (kernel linear, C=1)
[[402   5]
 [  5 38]]
```

## 21 - Matriz de confusão - Regressão Logística

```
In [29]: from sklearn.linear_model import LogisticRegression

lr = LogisticRegression().fit(X_train, y_train)
lr_predicted = lr.predict(X_test)
confusion = confusion_matrix(y_test, lr_predicted)

print('Regressão Logística\n', confusion)
```

```
Regressão Logística
[[401   6]
 [  8 35]]

/Users/marinaramalhetedesouza/opt/anaconda3/envs/ml-imp/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
n\_iter\_i = \_check\_optimize\_result(

## 22 - Matriz de confusão - Árvore de decisão

```
In [30]: from sklearn.tree import DecisionTreeClassifier

dt = DecisionTreeClassifier(max_depth=3).fit(X_train, y_train)
tree_predicted = dt.predict(X_test)
confusion = confusion_matrix(y_test, tree_predicted)

print('Árvore de decisão (max_depth = 2)\n', confusion)
```

```
Árvore de decisão (max_depth = 2)
[[404   3]
 [ 13 30]]
```



# Métricas de avaliação para classificação binária

Acurácia =  $TP + TN / (TP + TN + FP + FN)$

Precisão =  $TP / (TP + FP)$  aka PPV (Positive predictive value)

Recall =  $TP / (TP + FN)$  aka TPR (True Positive Rate)

$F1 = 2 \text{ Precisão Recall} / (\text{Precisão} + \text{Recall})$

## 23 - Computando métricas

```
In [32]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

print('Acurácia: {:.2f}'.format(accuracy_score(y_test, tree_predicted)))
print('Precisão: {:.2f}'.format(precision_score(y_test, tree_predicted)))
print('Recall: {:.2f}'.format(recall_score(y_test, tree_predicted)))
print('F1: {:.2f}'.format(f1_score(y_test, tree_predicted)))
```

Acurácia: 0.96

Precisão: 0.91

Recall: 0.70

F1: 0.79

## 24 - Reporte combinado

```
In [33]: from sklearn.metrics import classification_report

print(classification_report(y_test, tree_predicted, target_names=['not 1', '1']))
```

	precision	recall	f1-score	support
not 1	0.97	0.99	0.98	407
1	0.91	0.70	0.79	43
accuracy			0.96	450
macro avg	0.94	0.85	0.89	450
weighted avg	0.96	0.96	0.96	450

## 25 - Outros reportes combinados

```
In [34]: print('Estratificado (dummy)\n',
          classification_report(y_test, y_classprop_predicted, target_names=['not 1', '1']))
print('SVM\n',
      classification_report(y_test, svm_predicted, target_names = ['not 1', '1']))
print('Regressão Logística\n',
      classification_report(y_test, lr_predicted, target_names = ['not 1', '1']))
print('Árvore de Decisão\n',
      classification_report(y_test, tree_predicted, target_names = ['not 1', '1']))
```

Estratificado (dummy)	precision	recall	f1-score	support
not 1	0.90	0.91	0.91	407
1	0.10	0.09	0.10	43
accuracy			0.83	450
macro avg	0.50	0.50	0.50	450
weighted avg	0.83	0.83	0.83	450

SVM					
		precision	recall	f1-score	support
	not 1	0.99	0.99	0.99	407
	1	0.88	0.88	0.88	43
	accuracy			0.98	450
	macro avg	0.94	0.94	0.94	450
	weighted avg	0.98	0.98	0.98	450

Regressão Logística					
		precision	recall	f1-score	support
	not 1	0.98	0.99	0.98	407
	1	0.85	0.81	0.83	43
	accuracy			0.97	450
	macro avg	0.92	0.90	0.91	450
	weighted avg	0.97	0.97	0.97	450

Árvore de Decisão					
		precision	recall	f1-score	support
	not 1	0.97	0.99	0.98	407
	1	0.91	0.70	0.79	43
	accuracy			0.96	450
	macro avg	0.94	0.85	0.89	450
	weighted avg	0.96	0.96	0.96	450

## Funções de Decisão

### 26 - Função de decisão

```
In [35]: X_train, X_test, y_train, y_test = train_test_split(X, y_binary_imbalanced, random_state=0)
y_scores_lr = lr.fit(X_train, y_train).decision_function(X_test)
y_score_list = list(zip(y_test[0:20], y_scores_lr[0:20]))

y_score_list
```

```
/Users/marinaramalhetedesouza/opt/anaconda3/envs/ml-imp/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
Out[35]: [(0, -29.82878885428698),
(0, -19.38292791420195),
(0, -29.198327257201743),
(0, -21.746174115798574),
(0, -22.64239804996128),
(0, -11.80601316319518),
(1, 6.496016572624416),
(0, -23.35456282787909),
(0, -27.543436971220387),
(0, -26.88821968799928),
(0, -31.86269074250119),
(0, -22.486131395524808),
(0, -25.31799892081813),
```

```
(0, -13.384564231087923),
(0, -13.565608315834),
(0, -13.308404562543435),
(1, 12.180778621399947),
(0, -34.36249371177895),
(0, -13.231503124195685),
(0, -29.593934459364586)]
```

## 27 - Função Proba

In [36]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y_binary_imbalanced, random_state=0)
y_proba_lr = lr.fit(X_train, y_train).predict_proba(X_test)
y_proba_list = list(zip(y_test[0:20], y_proba_lr[0:20,1]))

y_proba_list
```

```
/Users/marinaramalhetedesouza/opt/anaconda3/envs/ml-imp/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

Out[36]:

```
n_iter_i = _check_optimize_result(
[(0, 1.1105077627365785e-13),
(0, 3.8203342505101854e-09),
(0, 2.0860638224426631e-13),
(0, 3.5954775289931444e-10),
(0, 1.4673424093995484e-10),
(0, 7.459511542756899e-06),
(1, 0.9984928349973813),
(0, 7.198503772008908e-11),
(0, 1.0915341073965711e-12),
(0, 2.1018153589129763e-12),
(0, 1.4528113762683703e-14),
(0, 1.715525648508253e-10),
(0, 1.0104917586429055e-11),
(0, 1.5387105925755556e-06),
(0, 1.2838982508367566e-06),
(0, 1.6604760443198858e-06),
(1, 0.9999948719431359),
(0, 1.1927751884236894e-15),
(0, 1.793207006513804e-06),
(0, 1.404486491577931e-13)]
```

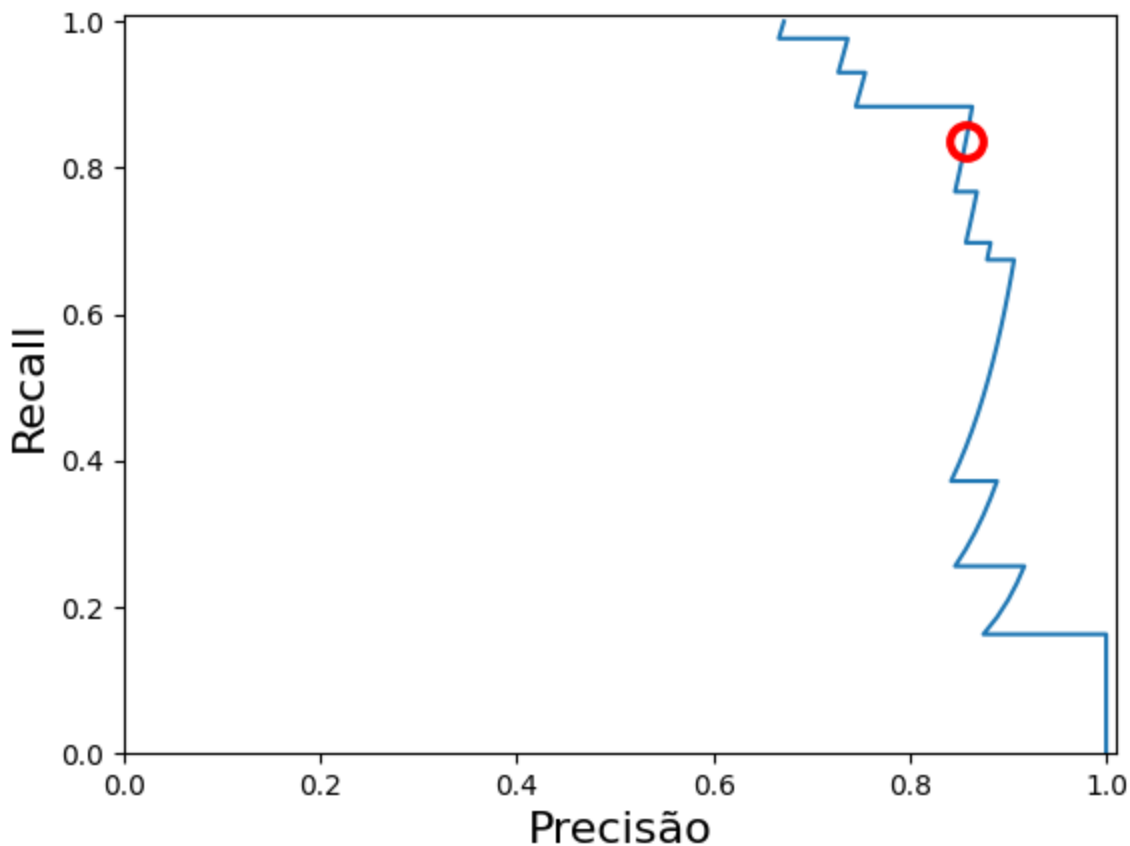
## 28 - Curvas de Precisão-Recall

In [46]:

```
from sklearn.metrics import precision_recall_curve

precision, recall, thresholds = precision_recall_curve(y_test, y_scores_lr)
closest_zero = np.argmax(np.abs(thresholds))
closest_zero_p = precision[closest_zero]
closest_zero_r = recall[closest_zero]

plt.figure()
plt.xlim([0.0, 1.01])
plt.ylim([0.0, 1.01])
plt.plot(precision, recall, label='Curva de Precisão-Recall')
plt.plot(closest_zero_p, closest_zero_r, 'o', markersize = 12, fillstyle = 'none', c='r',
plt.xlabel('Precisão', fontsize=16)
plt.ylabel('Recall', fontsize=16)
# plt.axes().set_aspect('equal')
plt.show()
```



## 29 - Área abaixo da curva ROC (Característica de Operação do Receptor)

In [48]:

```
from sklearn.metrics import roc_curve, auc

X_train, X_test, y_train, y_test = train_test_split(X, y_binary_imbalanced, random_state=0)

y_score_lr = lr.fit(X_train, y_train).decision_function(X_test)
fpr_lr, tpr_lr, _ = roc_curve(y_test, y_score_lr)
roc_auc_lr = auc(fpr_lr, tpr_lr)

plt.figure()
plt.xlim([-0.01, 1.00])
plt.ylim([-0.01, 1.01])
plt.plot(fpr_lr, tpr_lr, lw=3, label='Regressão Logística (area = {:.2f})'.format(roc_auc_lr))
plt.xlabel('Taxa FP', fontsize=16)
plt.ylabel('Taxa TP', fontsize=16)
plt.title('Curva ROC (classificador de dígitos)', fontsize=16)
plt.legend(loc='lower right', fontsize=13)
plt.plot([0, 1], [0, 1], color='navy', lw=3, linestyle='--')
# plt.axes().set_aspect('equal')
plt.show()
```

/Users/marinaramalhetedesouza/opt/anaconda3/envs/ml-imp/anaconda3/lib/python3.8/site-packages/sklearn/linear\_model/\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1): STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.


Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```



### 30 - Curvas ROC

In [49]:

```
from matplotlib import cm

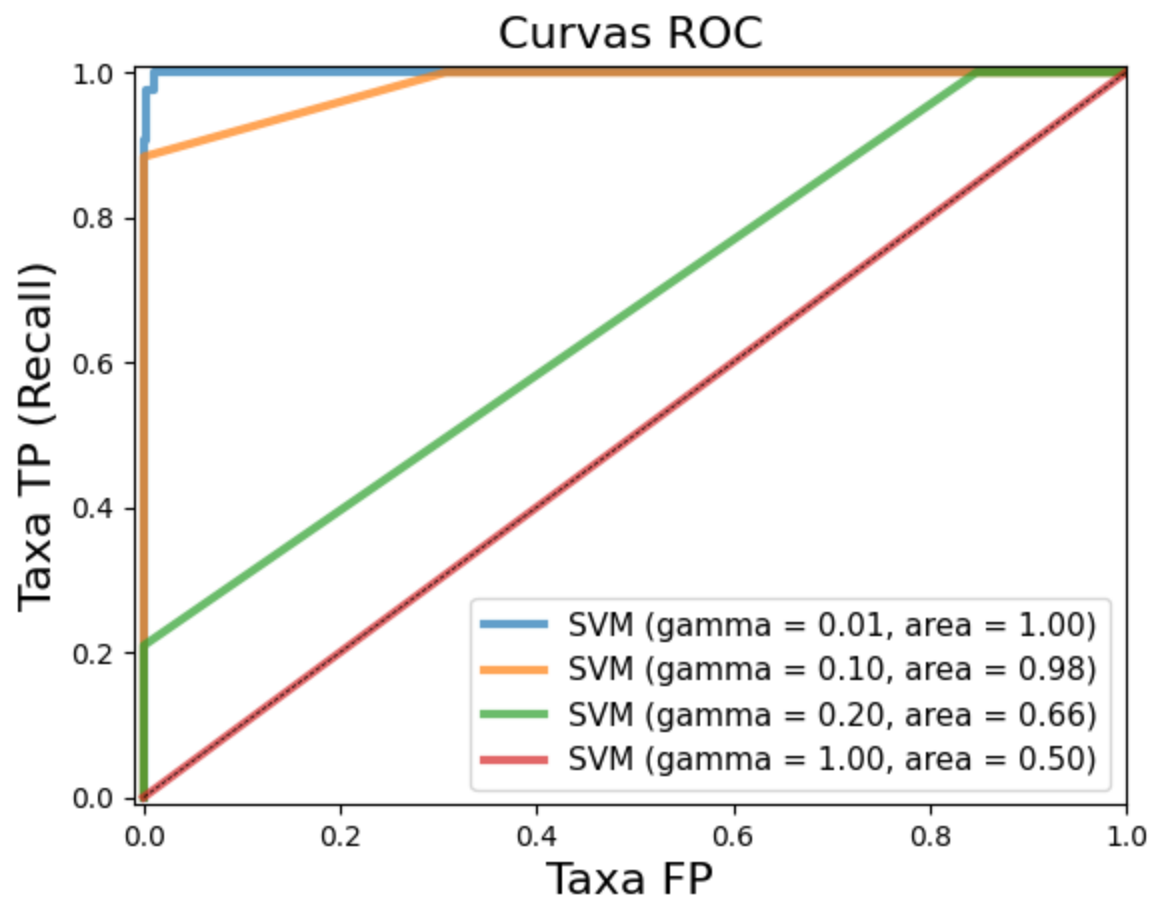
X_train, X_test, y_train, y_test = train_test_split(X, y_binary_imbalanced, random_state=0)

plt.figure()
plt.xlim([-0.01, 1.00])
plt.ylim([-0.01, 1.01])
for g in [0.01, 0.1, 0.20, 1]:
    svm = SVC(gamma=g).fit(X_train, y_train)
    y_score_svm = svm.decision_function(X_test)
    fpr_svm, tpr_svm, _ = roc_curve(y_test, y_score_svm)
    roc_auc_svm = auc(fpr_svm, tpr_svm)
    accuracy_svm = svm.score(X_test, y_test)
    print("gamma = {:.2f}  accuracy = {:.2f}  AUC = {:.2f}".format(g, accuracy_svm,
                                                                    roc_auc_svm))

    plt.plot(fpr_svm, tpr_svm, lw=3, alpha=0.7,
             label='SVM (gamma = {:.2f}, area = {:.2f})'.format(g, roc_auc_svm))

plt.xlabel('Taxa FP', fontsize=16)
plt.ylabel('Taxa TP (Recall)', fontsize=16)
plt.plot([0, 1], [0, 1], color='k', lw=0.5, linestyle='--')
plt.legend(loc="lower right", fontsize=11)
plt.title('Curvas ROC', fontsize=16)
# plt.axes().set_aspect('equal')

plt.show()
```



gamma = 0.01	accuracy = 0.91	AUC = 1.00
gamma = 0.10	accuracy = 0.90	AUC = 0.98
gamma = 0.20	accuracy = 0.90	AUC = 0.66
gamma = 1.00	accuracy = 0.90	AUC = 0.50

In [ ]: